

# SENTIMENT ANALYSIS OF OPAY APP REVIEWS

- BY: GBADAMOSI OLUWASEYI EMMANUEL
- FINAL YEAR PROJECT

# INTRODUCTION

Sentiment analysis, is an approach to natural language processing (NLP) that identifies the emotional tone behind a body of text. This is a popular way for organizations to determine and categorize opinions about a product, service or idea. It involves the use of data mining, machine learning (ML) and artificial intelligence (AI) to mine text for sentiment and subjective information. In recent years, several methods of sentiment analysis have been developed for many domains such as the health sector both in terms of disease and health services, entertainment fields such as film reviews, music, to the political field.



## AIM AND OBJECTIVES OF THE STUDY

The aim of this project is to analyze the sentiments about opay applications in the Google Play Store and to determine the level of user satisfaction based on comments in the Google Play Store review column and to determine the level of service success, shortcomings and how to further improve the app for better customer experience based on community sentiments. The objectives include:

1. Classifying the extracted data using Random forest as a supervised data mining algorithm to predict good and bad reviews.
2. To analyzing the reviews of users on opay app to extract sentiment
3. To get valuable insight on the level of app success, shortcomings and how to further improve the app for better customer experience

## RESULT SUMMARY

The trained random forest classification model had an accuracy of 86%, precision of 87%, recall of 86% and average cross validated ROC AUC of 90% indicating a relatively good model. The most important features were the 4 sentiment scores (positive, compound, neutral and negative), word and character counts, a few document vectors, and the words great and easy (from TD-IDF). The model was then applied to the neutral reviews to try and categorize the review into positive or negative.

- The RFC model predicted 564 reviews as good incorrectly and 445 review as bad incorrectly. The PR curve shows the calculated precision and recall at various threshold values. The precision values for our model remain relatively stable at each threshold AP= 0.95
- After extracting sentiment from the reviews using vader module, positive reviews had the highest number of 73% while negative reviews had the lowest number of 27%. According to the number of reviews positive reviews had the highest number of reviews which means people are satisfied with the app but we can't base the app success on only the positive reviews we have to strike a balance in order to get good insight,
- App success: Result shows that users are happy and like the app because it is fast, efficient and easy to use. Some key functions that users enjoyed were Oride service, electricity bill payment, airtime fill up, TV subscriptions amidst other services.
- Shortcoming: many of users complained of experiencing technical issues of the app not able to connect to the internet on their mobile phones, also geo-location for oride service not accurate causing late arrival of riders , riders misbehave, agents Unable get commission on

```
In [1]: import warnings
warnings.filterwarnings('always')
warnings.filterwarnings('ignore')
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.metrics import accuracy_score, classification_report, confusion_m
atrix, roc_curve, auc
from sklearn.metrics import roc_auc_score, average_precision_score, precision_
recall_curve
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score
# from sklearn.utils.fixes import signature
from nltk.corpus import wordnet, stopwords
from nltk import pos_tag, ngrams
from nltk.tokenize import WhitespaceTokenizer
from nltk.stem import WordNetLemmatizer
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from nltk.collocations import * #N-grams
import nltk.collocations #N-grams
from collections import Counter #N-grams
from gensim.test.utils import common_texts
from gensim.models.doc2vec import Doc2Vec, TaggedDocument
from wordcloud import WordCloud
import string
import seaborn as sns
import matplotlib.pyplot as plt; plt.rcParamsDefaults()
from matplotlib import rc
import missingno as msno #missmap
import numpy as np
import pandas as pd
import os
```

```
In [50]: #Set home drive
os.chdir("C:/Users/Oluwaseyi/Documents/final year project code/sentiment-analy
sis-on-opay-app/")#set drive
pd.set_option('display.max_colwidth', -1) #set column width for better string
viewing
```

## Exploratory Data Analysis (EDA)

This data was provided by a third party Appfollow. To perform similar analysis, a sample data set obtained from appfollow website.

## Data Description

The app store data used here contains 44,490 observations and 26 columns. Each customer review is composed of Date, AppID, AppName, Language, Version, Version Code, Rating, Title, Review, Translated title, Translated review, Reply Date, Developer Reply, User, Device, Device Type, Tags, Categories, Notes, Likes, Dislikes, Link, Permalink, AF Link.


```
In [51]: #Loading the data and initial preview
df = pd.read_csv("reviews_googleplay_1591435589.csv", encoding = "ISO-8859-1")
#Loading the data
print('Dimensions:',df.shape) #call data dimensions
df.dtypes
df.head()
```

Dimensions: (44490, 26)

Out[51]:

	Date	AppID	AppName	Language	Version	VersionCode	OS	Author	Rating
0	6/4/2018 6:48	team.opay.pay	OPay - OMall, ORide, Airtime, Transfer & more	en	NaN	NaN	NaN	Ethel Mwanyalo	5
1	6/11/2018 10:35	team.opay.pay	OPay - OMall, ORide, Airtime, Transfer & more	en	NaN	NaN	NaN	Abiola Daramola	5
2	6/13/2018 12:06	team.opay.pay	OPay - OMall, ORide, Airtime, Transfer & more	en	NaN	NaN	NaN	Yahaya Saminu	5
3	6/15/2018 11:28	team.opay.pay	OPay - OMall, ORide, Airtime, Transfer & more	en	NaN	NaN	NaN	A Google user	5
4	6/18/2018 11:40	team.opay.pay	OPay - OMall, ORide, Airtime, Transfer & more	en	NaN	NaN	NaN	Abena Music	5

5 rows × 26 columns




## Data Integrity

We can see that the fields TranslatedTitle, TranslatedReview, User, Device, DeviceType, Tags, Version, Version Code, os, Title, Categories and Notes have no data. These can be removed during the data cleaning phase. Additionally, ReplyDate, DeveloperReply, Title and Author are quite sparsely populated fields.

```
In [52]: msno.bar(df)
```

```
Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x45843a88>
```



## Initial Data Cleaning

To simplify the data, features that are less meaningful to our analysis or are too scarcely populated are removed from the dataset.

Author, User – Removed unique identifiers.

AppName, AppID - is in a cleaner format and more consistent. but of no meaning to the analysis

.

Language – The reviews should all be in English so this field is not meaningful

TranslatedTitle, TranslatedReview, User, Device, DeviceType, ReplyDate, DeveloperReply, Tags, Notes, Version, Version Code, os, Title, Categories – too scarcely populated or have no data to provide any meaningful insight.

Link – Removed as outside of the scope of this project, maybe a field of interest for further analysis.

.

```
In [53]: df = df.drop(["Author", "AppName", "Version", "VersionCode", "OS", "Language", "Translated title", "Translated review", "Reply Date", "Developer Reply", "User", "Device", "Device Type", "Tags", "Notes", "Link", "Likes", "Dislikes", "Permalink", "AF Link", "Categories", "Title"], axis=1)
df['AppID'] = df['AppID'].replace({'team.opay.pay': 'Opay'})
print('Dimensions:', df.shape)
df.head()
```

Dimensions: (44490, 4)

Out[53]:

	Date	AppID	Rating	Review
0	6/4/2018 6:48	Opay	5	Great stuff loading
1	6/11/2018 10:35	Opay	5	Good
2	6/13/2018 12:06	Opay	5	Nice
3	6/15/2018 11:28	Opay	5	Cool app
4	6/18/2018 11:40	Opay	5	How do it work to earn..pls urgently..

```
In [54]: #Total number rows and columns
df.shape
```

Out[54]: (44490, 4)

```
In [55]: # remove duplicates/ for every duplicate we will keep only one row of that type.
df.drop_duplicates(subset=['Rating', 'Review'], keep='first', inplace=True)
print(df.shape)
df.head()
```

(26535, 4)

Out[55]:

	Date	AppID	Rating	Review
0	6/4/2018 6:48	Opay	5	Great stuff loading
1	6/11/2018 10:35	Opay	5	Good
2	6/13/2018 12:06	Opay	5	Nice
3	6/15/2018 11:28	Opay	5	Cool app
4	6/18/2018 11:40	Opay	5	How do it work to earn..pls urgently..

```
In [56]: df["Reviews"]=df["Review"]
df.drop(['Review'],axis=1,inplace=True)
df.head()
```

Out[56]:

	Date	AppID	Rating	Reviews
0	6/4/2018 6:48	Opay	5	Great stuff loading
1	6/11/2018 10:35	Opay	5	Good
2	6/13/2018 12:06	Opay	5	Nice
3	6/15/2018 11:28	Opay	5	Cool app
4	6/18/2018 11:40	Opay	5	How do it work to earn..pls urgently..

## DATA PREPARATION AND CLEANING

My initial attempts at lemmatizing the review text were unsuccessful as a spot check of the corpus showed many words that were not transformed to their base form. Upon further research, it was noted that the default setting for the lemmatization module in NLTK wordnet was 'noun' resulting in the transformation of only noun words. To resolve this, the function below defines the word type based on the position tag obtained from the NLTK pos\_tag module (the pos\_tag module is applied in the clean\_text function in the following section)

```
In [57]: def get_tag(pos_tag):
    if pos_tag.startswith('J'):
        return wordnet.ADJ
    elif pos_tag.startswith('V'):
        return wordnet.VERB
    elif pos_tag.startswith('N'):
        return wordnet.NOUN
    elif pos_tag.startswith('R'):
        return wordnet.ADV
    else:
        return wordnet.NOUN
```

The clean\_text function defined below applies the following transformations:

- 1) Change all words to lower case (lemmatization does not work on capitals as they are assumed to be proper nouns).
- 2) Tokenize the text and remove punctuation.
- 3) Remove numeric values.
- 4) Remove stop words (using pre-built stop word dictionary).
- 5) Remove any empty tokens.
- 6) Apply a position tag to each word and define it based on the previously defined get\_tag function as adjective, noun, verb, or adverb.
- 7) Lemmatize the words.
- 8) Remove any single letter words resulting from lemmatization.



```
In [58]: import re
def clean_text(text):
    text = text.lower() #change all text to lower case
    text = re.sub(r'\s*(?:https?:\/\/\/?[\w.-]+(?:\.[\w.-]+)+[\w\-\._~:/?#[\]@!\$&\'\"(\)\*\+,;=.]+)', '', text)
    text = [word.strip(string.punctuation) for word in text.split(" ")] #token
    #ize and remove punctuation
    text = [word for word in text if not any(c.isdigit() for c in word)] #remo
    #ve numeric values
    stop = stopwords.words('english') #call english stop word dictionary
    text = [x for x in text if x not in stop] #remove stop words
    text = [t for t in text if len(t) > 0] #remove empty tokens
    pos_tags = pos_tag(text) #apply position tag to text
    text = [WordNetLemmatizer().lemmatize(t[0], get_tag(t[1])) for t in pos_ta
    #gs] #apply pos_tag function and lemmatize text
    text = [t for t in text if len(t) > 1] # remove single letter words
    text = " ".join(text) #combine
    return(text)
#create new column with cleaned text
df["reviews_clean"] = df["Reviews"].apply(lambda x: clean_text(x))
```

```
In [59]: #Text Before Text Cleaning
print('Before Text Cleaning')
df['Reviews'].head()
```

Before Text Cleaning

```
Out[59]: 0    Great stuff loading
1    Good
2    Nice
3    Cool app
4    How do it work to earn..pls urgently..
Name: Reviews, dtype: object
```

```
In [60]: #Text Before Text Cleaning
print('After Text Cleaning')
df['reviews_clean'].head()
```

After Text Cleaning

```
Out[60]: 0    great stuff loading
1    good
2    nice
3    cool app
4    work urgently
Name: reviews_clean, dtype: object
```


```
In [61]: # Drop all columns that are blank as a results of the text cleaning function.
         # Lost 176 rows.
print(df.shape)
df = df[df['reviews_clean'].map(len) > 0]
print(df.shape)
```

```
(26535, 5)
(26359, 5)
```

## Plot: Distribution of Reviews by Rating

This plot views the distributions of reviews across all ratings. We can see that the number of positive reviews has the highest number of reviews, causing our dataset to be imbalanced, followed by negative reviews having less than 6000 comments.

```
In [62]: sns.set(style="darkgrid")
bx = sns.countplot(x = "Rating", data=df)
bx.set(xlabel='Rating', ylabel='Number of Reviews',title='Distribution of Reviews by Rating')
plt.show()
```



## Most Frequently Occuring Words

```
In [63]: from sklearn.feature_extraction.text import CountVectorizer

cv = CountVectorizer(stop_words = 'english')
words= cv.fit_transform(df['reviews_clean'].values.astype('U'))

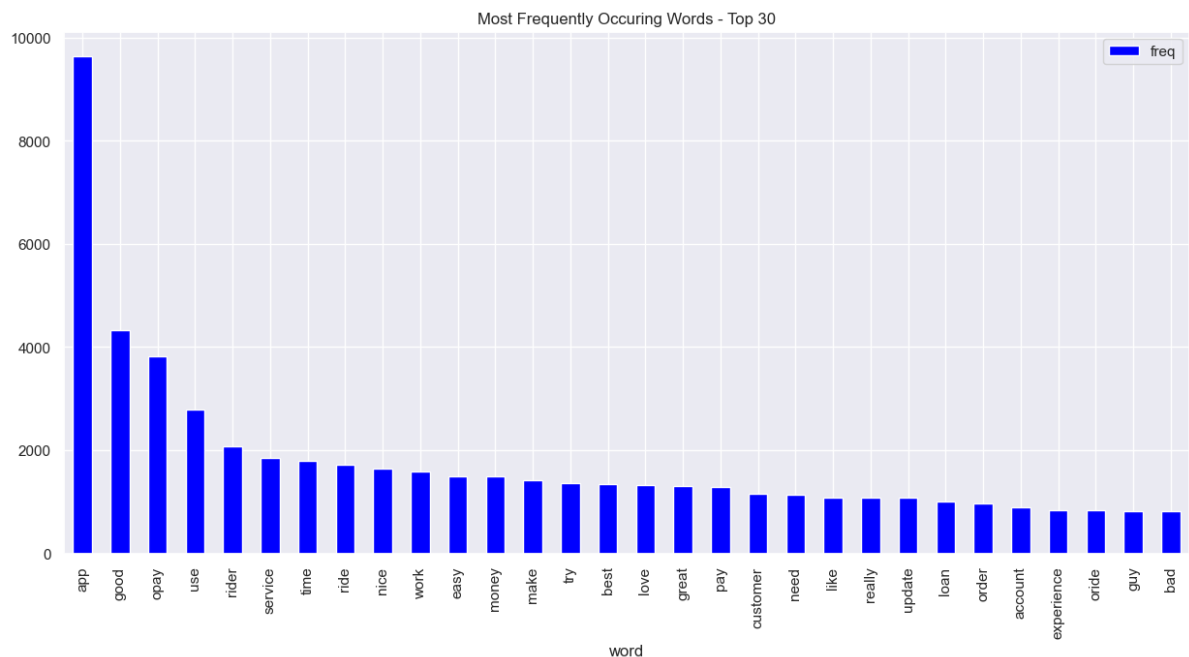
sum_words = words.sum(axis=0)

words_freq = [(word, sum_words[0, i]) for word, i in cv.vocabulary_.items()]
words_freq = sorted(words_freq, key = lambda x: x[1], reverse = True)

frequency = pd.DataFrame(words_freq, columns=['word', 'freq'])

frequency.head(30).plot(x='word', y='freq', kind='bar', figsize=(15, 7), color = 'blue')
plt.title("Most Frequently Occuring Words - Top 30")
```

Out[63]: Text(0.5, 1.0, 'Most Frequently Occuring Words - Top 30')



```
In [64]: dfn= df[df['reviews_clean'].str.contains("app")]
#dfn_pg['Reviews'].head(10)
dfn['reviews_clean'].head(10)
```

```
Out[64]: 3      cool app
6      apply loan
11     useless ewallet balance even reflect app can't buy \nairtime pay bill
i'm fear money pump name \nof top ewallet
14     least let get update work tire update app \ndoesn't get transaction exe
cute
16     think best app e-wallet transfer work perfectly \nin mine
32     best app sow far please need customer phone number
35     good app credit debit also use buy credit data
36     can't send money bank account use app top account cant send money bad
40     good app user friendly
43     best money app
Name: reviews_clean, dtype: object
```

## FEATURE ENGINEERING

### Sentiment Analysis

The Vader module from NLTK was the model selected for sentiment analysis. The Vader module uses a prebuilt lexicon of words to calculate a sentiment score. This module was selected for sentiment analysis because the module takes into consideration the context of the text. The module returns 4 values: positivity score, neutrality score, negativity score and summary score.


```
In [65]: sid = SentimentIntensityAnalyzer()
#calculates the negativity, neutrality, positivity and overall sentiment scores
df["sentiments"] = df["reviews_clean"].apply(lambda x: sid.polarity_scores(x))
#drop sentiment column and add the 4 sentiment scores as separate features to primary dataset
df = pd.concat([df.drop(['sentiments'], axis=1), df['sentiments'].apply(pd.Series)], axis=1)
df[['AppID', 'Rating', 'reviews_clean', 'neg', 'neu', 'pos', 'compound']].head(10)
```

Out[65]:

	AppID	Rating	reviews_clean	neg	neu	pos	compound
0	Opay	5	great stuff loading	0.0	0.328	0.672	0.6249
1	Opay	5	good	0.0	0.000	1.000	0.4404
2	Opay	5	nice	0.0	0.000	1.000	0.4215
3	Opay	5	cool app	0.0	0.303	0.697	0.3182
4	Opay	5	work urgently	0.0	1.000	0.000	0.0000
5	Opay	2	unable either send receive cash seem like fully \nunsupported	0.0	0.579	0.421	0.6240
6	Opay	5	apply loan	0.0	1.000	0.000	0.0000
7	Opay	1	can't get access code verify phone number	0.0	0.822	0.178	0.0772
8	Opay	4	love	0.0	0.000	1.000	0.6369
9	Opay	4	can't find option put amount want recharge	0.0	0.822	0.178	0.0772

The graph below shows the compound sentiment calculated by Vader distributed by good and bad reviews. We can see that good reviews are mostly considered very positive by Vader, whereas, bad reviews are more dispersed with a slightly higher proportion of negative reviews with negative sentiment scores. The only variation to this trend is the slight peak around the neutral compound score (zero) for both negative and positive reviews.

```
In [66]: for x in [1, 5]:
        subset = df[df['Rating'] == x]
        if x > 3:
            label = "Postive Review"
        else:
            label = "Negative Review"
        sns.distplot(subset['compound'], hist = False, label = label).set_title('S
        entiment Score Distribution by Review Type')
```



## Word and Character Count Features

Two new features are created by extracting the number of characters and number of words per review. Log transformation is applied to pull in outliers.


```
In [67]: df["num_chars"] = df["reviews_clean"].apply(lambda x: len(x))
        df["num_words"] = df["reviews_clean"].apply(lambda x: len(x.split(" ")))
        #Log transformation
        df['num_chars1'] = np.log(df['num_chars'])
        df['num_words1'] = np.log(df['num_words'])
```

```

In [68]: x1 = df['Rating']
x2 = df['Rating']
y1 = df['num_chars1']
y2 = df['num_words1']
#plot num_chars by rating in column 1
plt.subplot(1, 2, 1)
plt.xticks(rotation=90)
g = sns.boxplot(x1, y1)
g.set(xlabel='Rating', ylabel='Number of Characters',title='')
#plot num_words by rating in column 2
plt.subplot(1, 2, 2)
g = sns.boxplot(x2, y2)
g.set(xlabel='Rating', ylabel='Number of Words',title='')

plt.tight_layout()
plt.show()

```



We can see a trend forming from the boxplot below, where users tend to leave longer reviews for negative ratings ( $<3$ ) and neutral ratings ( $=3$ ) and shorter reviews good review ( $>3$ ). This may be a useful feature for our predictive models

## Doc2Vec Feature Creation

The doc2vec method from the Genism module is used to generate document vectors for each cleaned review. The doc2vec module uses a modified word2vec model with the addition of a document unique vector, which numerically represents the document. This provides a document-concept representation of each review. This feature is important for training our model since similar texts should have similar vector representations. We first start by creating doc2vec vector columns and then proceed to train the model. The model is then applied to the text to transform each review into vector data before being combined with our original dataframe.

Warning Message to install compiler to speed up genism is not necessary for the size of data used in this notebook. For larger data, a compiler would be recommended as this model took roughly 26 minutes to run.

```
In [69]: documents = [TaggedDocument(doc, [i]) for i, doc in enumerate(df["reviews_clean"].apply(lambda x: x.split(" ")))]
# train a Doc2Vec model with our text data
model = Doc2Vec(documents, vector_size=6, window=2, min_count=1, workers=4)
# transform each document into a vector data
df_vector = df["reviews_clean"].apply(lambda x: model.infer_vector(x.split(" "))).apply(pd.Series)
df_vector.columns = ["df_vector_" + str(x) for x in df_vector.columns]
df = pd.concat([df, df_vector], axis=1)
```

## Term Frequency - Inverse Document Frequency

The word frequency is calculated using the TF-IDF model. In addition to just counting word frequency, this model computes the relative importance of each word based on the frequency of occurrence of the word in each text. A column is generated for every word which occurs in a minimum of 10 different documents to provide a relative filter on importance and to remove size. This can be adjusted to fine tune the predictive models.

```
In [70]: # add tf-idfs columns
from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(min_df = 10)
tfidf_result = tfidf.fit_transform(df["reviews_clean"]).toarray()
tfidf_df = pd.DataFrame(tfidf_result, columns = tfidf.get_feature_names())
tfidf_df.columns = ["word_" + str(x) for x in tfidf_df.columns]
tfidf_df.index = df.index
df = pd.concat([df, tfidf_df], axis=1)
```



## Define Good and Bad Reviews

The final feature created is to define a bad review ( rating < 3) by denoting it with 0 and all other ratings with 1. For the purposes of our model, the neutral reviews (rating of 3) are separated into another dataframe. Our dataset is relatively imbalanced with 73% good review and 27.0% bad reviews oversampling of our dataset would take care of the imbalanced to make it a balanced dataset.

In [71]: `df.shape`

Out[71]: (26359, 1562)


```
In [72]: df['label'] = np.where(df['Rating']<3, 0, 1)
#take lowest and highest rating
df_class = df[(df['Rating'] < 3) | (df['Rating'] > 3)]
df_neutral = df[(df['Rating'] == 3)].drop(['label'], axis=1)
df_class = df_class.sort_values(by=['Rating'])
print ("Dimenions:", df_class.shape)
print ("Good (1) vs Bad (0) split:" "\n",df_class["label"].value_counts(normal
ize = True))
df_class.groupby('label').count()
```

```
Dimenions: (24142, 1563)
Good (1) vs Bad (0) split:
1    0.729268
0    0.270732
Name: label, dtype: float64
```

Out[72]:

	Date	AppID	Rating	Reviews	reviews_clean	neg	neu	pos	compound	num_ch
label										
0	6536	6536	6536	6536	6536	6536	6536	6536	6536	6536
1	17606	17606	17606	17606	17606	17606	17606	17606	17606	17606

2 rows × 1562 columns



## MODEL DEVELOPMENT

The Random Forest model (RF) is used to predict if a review is good or bad given the various features we created from the review text. The model will then be used on the neutral dataset (rating = 3) to categorize the reviews.

### Random Forest Classifier

The features used to train the RF model are selected and any columns to be ignored are defined. The dataset is then split into training and test datasets.

```
In [73]: # feature selection
label = "label"
ignore_cols = [label, "Reviews", "reviews_clean", "Date", "AppID", "Rating"]
features = [c for c in df_class.columns if c not in ignore_cols]
# split the data into train and test
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df_class[features], df_class[label], test_size = 0.3, random_state = 42)
```

The resultant dataset for training is 16,899 rows x 1,558 columns and the test dataset is 7,243 rows x 1,558 columns. The 80/30 split was used as the app dataset is relatively small.

```
In [74]: print('Training Features Shape:', X_train.shape)
print('Training Labels Shape:', y_train.shape)
print('Testing Features Shape:', X_test.shape)
print('Testing Labels Shape:', y_test.shape)
```

```
Training Features Shape: (16899, 1557)
Training Labels Shape: (16899,)
Testing Features Shape: (7243, 1557)
Testing Labels Shape: (7243,)
```

## Balancing our label class

```
In [75]: from imblearn.combine import SMOTETomek
```

```
In [76]: os=SMOTETomek(1)
X_train_ns,y_train_ns=os.fit_sample(X_train,y_train)
print("The number of classes before fit {}".format(Counter(y_train)))
print("The number of classes after fit {}".format(Counter(y_train_ns)))
```

```
The number of classes before fit Counter({1: 12319, 0: 4580})
The number of classes after fit Counter({1: 12157, 0: 12157})
```

imblearn is used to balance our label class using over sampling method so our model would not be biased when predicting and it would be able to predict well on new dataset. Our initial label class was 0: 12319, 1: 4580 after imblearn was used we had 0: 12312, 1: 12312 ratio 50:50 making it a balanced dataset.

The RF model is trained and cross validation is run to get a better overview of our model's performance..

```
In [77]: # train a random forest classifier
rf = RandomForestClassifier(n_estimators = 150, random_state = 101)
rf.fit(X_train_ns,y_train_ns)
#Cross Validation Score
rfc_cv_score = cross_val_score(rf, df_class[features], df_class[label], cv=10,
scoring= 'roc_auc')
```

we use pickle library to save our model to avoid retraining everytime we work on our dataset

```
In [97]: import pickle
with open('random_forest_classifier2.pickle','wb') as f:#saving our model
    pickle.dump(rf,f)

# pickle_in= open('random_forest_classifier1.pickle','rb')# opening our model
# clf= pickle.load(pickle_in)
```

## Model Evaluation

### Confusion Matrix:

The RFC model predicted 564 reviews as good incorrectly and 445 review as bad incorrectly.

### Classification Report:

The model achieved an average precision of 0.86, average recall of 0.86 and average accuracy of 0.86. We can see the model has higher precision when it comes to predicting positive review. This may be because the positive sentiment is one of most important feature for our model (refer to Feature Importance section).

### Cross Validated (CV) AUC Score:

The model achieved an average CV AUC score of 0.90 which indicates a relatively good model.

```
In [79]: print('CONFUSION MATRIX')
print(confusion_matrix(y_test, rf.predict(X_test)))
print('\n')
print(pd.crosstab(y_test, rf.predict(X_test), rownames=['Actual Result'], col
names=['Predicted Result']))
print('\n')
print('CLASSIFICATION REPORT')
print(classification_report(y_test, rf.predict(X_test)))
print('\n')
print('ALL AUC SCORES')
print(rfc_cv_score)
print('\n')
print('MEAN AUC SCORE: ', rfc_cv_score.mean())
```

CONFUSION MATRIX

```
[[1509  447]
 [ 544 4743]]
```

Predicted Result	0	1
Actual Result		
0	1509	447
1	544	4743

CLASSIFICATION REPORT

	precision	recall	f1-score	support
0	0.74	0.77	0.75	1956
1	0.91	0.90	0.91	5287
accuracy			0.86	7243
macro avg	0.82	0.83	0.83	7243
weighted avg	0.87	0.86	0.86	7243

ALL AUC SCORES

```
[0.86523547 0.89011491 0.90828208 0.90139439 0.91017906 0.93943955
 0.93123252 0.91447937 0.8901036 0.87808333]
```


MEAN AUC SCORE: 0.9028544261567095

## Receiver Operating Characteristics (ROC) Curve

The trade-off between the true positive (TP) and false positive (FP) rate is shown in the Receiver Operating Characteristics (ROC) curve, and can be used to assess the quality of the classifier used in our model. The distance between the ROC curve and the diagonal baseline indicates the reliability of the predictions from our model. The model is quite good with an area under curve (AUC) value of 0.91.

Note: ROC is not a good indicator of model quality if the data is skewed towards a specific outcome as this could mute the FP and FN prediction rates (depending on the skewing of data). The app data review was relatively balanced in terms of the number of defined good or bad reviews, which give us some confidence in the ROC curve

```
In [80]: y_pred = [x[1] for x in rf.predict_proba(X_test)]
fpr, tpr, thresholds = roc_curve(y_test, y_pred, pos_label = 1)
roc_auc = auc(fpr, tpr)
plt.figure(1, figsize = (15, 10))
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()
```



## Precision Recall Curve (Average Precision)

The PR curve shows the calculated precision and recall at various threshold values. The precision values for our model remain relatively stable at each threshold  $AP = 0.95$ .

Precision (positive prediction value) is the ratio of  $TP / (TP + FP)$   
 Recall (sensitivity) is the ratio of  $TP / (TP + FN)$

Note: The PR curve is useful for dataset that are imbalanced.

```


In [81]: from funcsigs import signature
average_precision = average_precision_score(y_test, y_pred)
precision, recall, _ = precision_recall_curve(y_test, y_pred)
step_kwargs = ({'step': 'post'}
                if 'step' in signature(plt.fill_between).parameters
                else {})

plt.figure(1, figsize = (15, 10))
plt.step(recall, precision, color='b', alpha=0.2,
         where='post')
plt.fill_between(recall, precision, alpha=0.2, color='b', **step_kwargs)

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('2-class Precision-Recall curve: AP={0:0.2f}'.format(average_precision))

```

Out[81]: Text(0.5, 1.0, '2-class Precision-Recall curve: AP=0.95')



## Feature Importance

The most important features were the 4 sentiment scores generated by Vader, the doc2vec, the number of words and characters features. Additionally, some words identified by TF-IDF such as “app” and “good” have fairly high importance, also may be correlated with the Vader “pos” scores and the words identified by TF-IDF such as “bad” may be correlated with the Vader “neg” scores, while the words identified by TF-IDF such as “update”, “download”, “can” may be correlated with the Vader “neu” scores.

In [ ]:

```
In [82]: # show feature importance
feature_importances_df = pd.DataFrame({"feature": features, "importance": rf.feature_importances_}).sort_values("importance", ascending = False)
feature_importances_df.head(20)
```

Out[82]:

	feature	importance
2	pos	0.076610
3	compound	0.070698
0	neg	0.044265
1	neu	0.041633
9	df_vector_1	0.033498
7	num_words1	0.031259
5	num_words	0.030501
6	num_chars1	0.029769
8	df_vector_0	0.029394
12	df_vector_4	0.027643
4	num_chars	0.026245
10	df_vector_2	0.024661
13	df_vector_5	0.023368
11	df_vector_3	0.022032
103	word_app	0.014389
588	word_good	0.010922
149	word_bad	0.009752
1451	word_update	0.009485
399	word_download	0.009431
214	word_can	0.009093



## Model Application

The RF model is applied to the dataset with ratings of 3 to determine if the reviews are good or bad.

```
In [83]: df_temp = df_neutral[['Date', 'AppID', 'Rating', 'Reviews', 'reviews_clean']]
df_neutral = df_neutral.drop(['Reviews', 'reviews_clean', 'Date', 'AppID', 'Rating'], axis=1)
df_neutral['prediction'] = rf.predict(df_neutral)
df_neutral = pd.concat([df_temp.reset_index(drop=True), df_neutral.reset_index(drop=True)], axis=1)
```

## Preview Predicted Reviews

### Predicted Good Reviews

- The predicted good review preview seems to be less insightful, but this is somewhat expected given the high neutrality score noted during the Vader sentiment step. We see the model has failed to identify sarcasm in line 10: "Pay first before we rate the app", "please how can i upgrade my account? i want to fund more than N10,000 to my account" and "The app is great. But please trying including an option for transaction cancelation. And the airtime limit of 200 is much, the data too please add daily subscription too ...", "want to be an agent how can I go about it".
- In line 23: "The app is great. But please trying including an option for transaction cancelation. And the airtime limit of 200 is much, the data too please add daily subscription too ..." the app lacks this features rating may improve if in further update this feature are added to the app

We noted that our model weighed positive words heavily in feature importance, which would explain these results.

```
In [84]: print('Predicted Good Reviews')
dfn_pred_good = df_neutral[(df_neutral['prediction'] == 1)]
dfn_pred_good['Reviews'].head(20)
```

Predicted Good Reviews

```
Out[84]: 0      Good
2      Can I use these OPay and send money to Nigeria?
3      Its ik
6      Wonderful
7      Opay transfer charges are lesser than what banks charge,, and also the
fact that your pos transaction are processed same day is good. however your c
harge on transaction (0.99%) is quite high, also the number of weeks one has
to wait before he gets a machine should be treated. Again 6million per month
on transactions with machine is somehow high IMO.
8      It's been good so far but needs more perfection
9      good
10     perfect
12     Good App I like it, transactions that I've done on the app are just ins
tant. keep the good work and keep improving. I like to know how I can pay my
electricity bill on the app please can any of your amiable CR put me through
on that. thanks much.
15     Though the app is efficient and fast but so many features such as airti
me,transaction receipt printout,provision for saving an account as a benefici
ary,other betting companies,tv subcriptions etc. Thanks
16     please how can i upgrade my account? i want to fund more than N10,000 t
o my account
17     it's very good, there is room for improvement
18     best services
19     Nice to use
20     ok
22     want to be an agent how can I go about it
23     The app is great. But please trying including an option for transaction
cancelation. And the airtime limit of 200 is much, the data too please add da
ily subscription too
24     Mordred
27     fair
29     doing good
Name: Reviews, dtype: object
```

< >

## Predicted Bad Reviews

For the most part, the model seems to have done a pretty good job of categorizing the neutral reviews. Based on the preview we can see a couple of issues being highlighted by reviewers: : Alert notification issues : Sender unable to receive alert from Opay as there receipt

- Lack of services: use of debit card, POS you uable to receive money from other banks
- Error in services: Referral not working, unable to login, takes time to connect to the internet (app is slow), app unable to connect to mobile network, geolocation service not accurate
- Errors in transactions: failed EFT, failed deposit, etc
- Opay agents issues: Unable get commission on some services, charges on transcation is high
- Oride service issues: late arrival time of divers, divers are unmannered, divers reject long distance trip, drivers charge off the book

```
In [85]: print('Predicted bad Reviews')
dfn_pred_bad = df_neutral[(df_neutral['prediction'] == 0)]
dfn_pred_bad['Reviews'].head(50)
```

## Predicted bad Reviews

Out[85]: 1 Pay first before we rate the app

4 We want d sender to be receiving alert from Opay as there receipt dat  
d money have getting to d receiver account

5 The app is OK but my account was blocked and I have a lot of money in  
there

11 Pls i deposited some certain amount of money to my opay account from m  
y normal bank account it was deducted but opay showing my transaction failed  
why and no money deposited in my opay account

13 i was rub and my phone and sim was stolen how can i get access to my a  
ccount again cus i have money inside

14 is ok

21 Good app,but u need to give agents commission on some payments. Exampl  
es are electricity, data, cable TV etc. with that,we will be able to make mon  
ey on the platform.

25 Hi, I took a ride with ORide on the 30th of May 2019 from Allen, Ikeja  
to Mile 12, however, it appear that the ride was not ended by the rider and I  
was eventually charged over â€11000 after the rider driver end the ride the  
following day, please help fix it.

26 the app is very nice...but they should provide ATM card for withdrawer  
instead of transferring the money to bank account before you can withdraw mone  
y

28 ls good to know who cslded u becos some people can be stupid enough wh  
en u ask who is on the line they cut it off

31 Refuse to open

33 firsr ride was ok after i downloaded the app yesterday ,but since then  
if i request for ride i wait for hours and no bike shows up,.now the app is ju  
st telling me network failure ,..so not cool

34 i just set up my account and its already telling me that my account is  
logged in on another phone.

35 Oride is frustrating. It doesn't give request to nearby riders, but lo  
ng distance riders. Work on it. Thank you.

37 this service would have been great, except drivers now prefer to carry  
passengers off the app, rather than on the app. During rush hour in the eveni  
ng, you would see them line up in ikeja to pick up passengers who are not usi  
ng the app.

38 Urgh, I have to sign in every time, well that's not bad since it's als  
o a payment app, so security is paramount. I'm rating the app 3 cos I don't g  
et ride on time and I've been deducted double for one ride. Beyond this, the  
app is good

41 the app is working successful before but now it keep saying Something  
went wrong, check your internet connection, and my connection is good, and i  
transfer some amount and later it was refunded back to my bank account.

44 This app is wasting of time and was of ðŸ™¸

46 all the verification sent didnt get to my mail, instead i was blocked  
with my money inside...also the app geolocation needs to be properly checked,  
it doesn't send request to the nearest rider;rathet it sends it far away

47 please tell your riders to always pick customer request no matter the  
distance... and you people should make sure their papers are complete... beca  
use they keep saying they refuse some request because of the incomplete paper  
s to avoid being arrested... and harrassed... your promoters are the best

51 Bad App

54 Opay came to makurdi but refused to give me pos of which i created acc  
ount with them

55 ORide Is not a scam, Am having issue with pay out to the rider man tha  
t get me to my destination and Apart from maybe they are trying to resolve th  
e issue. Atleast yesterday i was able to recharge through it and my account g  
ot credited. So they should try to resolve the issue before i can give them 5

star thanks

56 not downloading

58 this app no longer open in my phone, every time I try to open it they say network error, but all other network using apps are working, why?, pls do something abt it

59 a bit delay on a request of a rider to pick u up at d locaton point

60 the app is nice and fast. but i made a transaction which failed and they're yet to reverse and refund my money. I've sent emails to them yet they've not reversed my money.

67 While the app is a very good development in terms of mobile money technology, but keep your agent for a particular period before entrusting them with the POS machine to me I not encouraging after asking for different document still that trust is not there, secondly without the POS you can't receive money from all the banks in Nigeria so how do you want your agent to coup and thirdly state mostly from the Northern part and North central can't top-up there electricity bills from this App i.e those using kaduna electricity company and PHCN. I'm registered as an agent no utility material, POS and am based in Kebbi State Nigeria.

71 Dear Opay, my App has refused to connect to internet for sometime now even though I have data unless connected to wifi, I have updated it, Uninstalled and re installed still the same, I hope you can help resolve this, and mostly riders are not available

72 I think you should work more on making your app user friendly. Then also work on high level professionalism from your drivers, some of them don't know how to talk to customers by insulting them and they always reject people's order or turn off their app and be looking for offline customers.

73 your referral system is not working and it took an hour 30 minutes plus to see a rider in ibadan

75 it good

76 i am just using it for the first time, made a payment through the mobile app it hasn't shown on my okash

77 why is the app not connecting to my internet even when i have a strong network

78 I have been using th App for days now without any problem. but today this App has been asking me to verify my phone number. even after inputting th OTP, I still can't login.

80 I am unable to access the app despite having enough internet data. This is unfortunate as I already made a deposit of 1,000 naira into the app for ORide. Pls admin what is happening?

81 I tried again and this time it was different. So 3 stars i thought you were having a promo?

82 It's difficult to get a ride. Even when drivers are in cluster close to me. Yet I can't still get a ride

83 This application can be sometimes frustrating in search for riders

84 The service is good though but you guys really need to work on our orders. It's annoying to see lots of riders and none of them would get our requests until someone afar picks and arrive late.

87 hard to get a bike

89 i don't really know what is happening with this app, i have been trying to download d app but it was not going thru

90 you cannot request for bikes even when they are right in front of you, they cannot start trips on their own and are difficult to assign most times

91 The site is telling me that the application will not be downloaded and installed on my phone Samsung galaxy tab sm-T561.

93 App refuses to connect on mobile connection except WiFi even after updating the app. I use an xiaomi brand

94 try to work on this new upgrade.. our order for bike don't always get r

esponse


96 i notice that u always cancelled a long distance trip but the moment i try to request for a short distance trip immediately a rider will be available. if u know the promo is for short distance trip, u should specify so that i will not waste my time requesting for ride.

97 The transport service is nice, the only issue is that they dont go to far distance, once you request for it, the riders wont connect, until the location is changed. Apart from that its fast and reliable.

100 location accuracy and finder, willingness of bike man to go far location and you really need much more advertisement.

101 I have been requesting for a ride since on Saturday last week 20/7/2019, for more than 30 mins, keep seeing non is available, it's so annoying, rather I requested for others & got it one hand.

Name: Reviews, dtype: object



## REVIEW INSIGHTS

Now that we have our cleaned review data and have split the neutral ratings into good or bad categories. We can examine the text to see what insights we can gather.

### Word Cloud

The word cloud is a visual representation of word frequency. We can immediately identify some key app services that seem to be important to customers, such as easy, good and opay. This method is somewhat controversial as it is difficult to interpret relative size (and therefore frequency) of words. It is also difficult to interpret context when isolated words are presented, such as in the case of "time, service, ride and work" which can be positive or negative.





```
In [87]: def n_gram(token, n_gram, size ):
          tokenized = token.apply(lambda x: x.split())
          finder = BigramCollocationFinder.from_documents(tokenized.values)
          bigram_measures = nltk.collocations.BigramAssocMeasures()
          finder.apply_freq_filter(1)
          result = finder.nbest(bigram_measures.pmi, 10)
          ngram_list = [pair for row in tokenized for pair in ngrams(row, n_gram)]
          counts = Counter(ngram_list).most_common()
          print (pd.DataFrame.from_records(counts, columns=['gram', 'count']).head(size))
```

Taking an initial look at the n-grams for the entire cleaned corpus, as we can see noise still exist in our clean corpus our model classified the noise as a positive phrase i guess the model mistook is as "okay" phrase. Apart from the noise, the positive phrase "good, app, great, easy" this can also denote that the app is good and it is easy to use

lets take a look at the negative N-gram reviews, we can see mostly negative phrases which make somewhat sense considering that negative reviews tend to have more text. The most prevalent complaints being Server error issues, app unable to connect to the internet and inaccurate geolocation service,. These are possible areas for the app developers to address to improve customer satisfaction.

## Positive N-Grams

Aside from the praise for the app, we can glimpse what customers like about opay app. the positive phrase "good, app, great, easy" this can also denote that the "app is good and it is easy to use". A successful app seems to be defined by the ability to make of it been easier, convenient and more accessible.

```
In [88]: df_best = df_class[(df_class['label'] == 1)]
          n_gram(df_best['reviews_clean'], 5, 15)
```

	gram	count
0	(ð???, ð???, ð???, ð???, ð???)	11
1	(one, best, app, ever, see)	4
2	(keep, tell, something, go, wrong)	4
3	(god, bless, opay, god, bless)	4
4	(best, app, i've, ever, use)	3
5	(ever, since, start, use, opay)	3
6	(great, app, easy, use, opay)	3
7	(always, tell, something, go, wrong)	2
8	(app, ever, keep, good, work)	2
9	(use, opay, make, life, easy)	2
10	(enroute, make, pay, service, render)	2
11	(nice, app, fast, payment, transaction)	2
12	(sometimes, take, long, get, rider)	2
13	(one, best, development, happen, nigeria)	2
14	(tell, set, high, accuracy, location)	2

## Negative N-Grams

lets take a look at the negative N-gram reviews, we can see mostly negative phrases which make somewhat sense considering that negative reviews tend to have more text. The most prevalent complaints being Server error issues, app unable to connect to the internet and inaccurate geolocation service.

```
In [89]: df_worst = df_class[(df_class['label'] == 0)]
         n_gram(df_worst['reviews_clean'], 5, 20)
```

	gram	count
0	(keep, say, something, go, wrong)	23
1	(keep, tell, something, go, wrong)	19
2	(something, go, wrong, please, check)	18
3	(go, wrong, please, check, connection)	12
4	(say, something, go, wrong, check)	11
5	(something, go, wrong, check, internet)	11
6	(something, go, wrong, check, connection)	11
7	(tell, something, go, wrong, check)	10
8	(wrong, please, check, connection, try)	10
9	(go, wrong, check, internet, connection)	9
10	(always, say, something, go, wrong)	9
11	(keep, show, something, go, wrong)	8
12	(please, check, connection, try, later)	7
13	(app, say, something, go, wrong)	6
14	(error, message, something, go, wrong)	6
15	(set, high, accuracy, location, service)	5
16	(use, app, say, something, go)	5
17	(app, keep, tell, something, go)	5
18	(something, go, wrong, check, network)	5
19	(app, keep, say, something, go)	5

## Predicted Positive N-gram

One interesting observation is the common occurrence of phrases like "good..", "easy..", etc. This seems to highlight some useful feedback for improving the apps and warrants a closer look.

```
In [90]: n_gram(dfn_pred_good['reviews_clean'], 3, 15)
```

	gram	count
0	(make, life, easy)	5
1	(keep, good, work)	3
2	(pay, electricity, bill)	3
3	(sometimes, take, time)	3
4	(still, need, improvement)	3
5	(rider, accept, ride)	3
6	(make, transportation, easy)	3
7	(nice, easy, use)	3
8	(app, great, service)	2
9	(great, experience, far)	2
10	(use, app, thank)	2
11	(app, good, easy)	2
12	(good, easy, use)	2
13	(good, one, fast)	2
14	(take, like, forever)	2

```
In [91]: dfn_pg= dfn_pred_good[dfn_pred_good['reviews_clean'].str.contains("easy")]
#dfn_pg['Reviews'].head(10)
dfn_pg['reviews_clean'].head(10)
```

```
Out[91]: 65      okay previous complaint resolve chaging rating sincerely like app try
        achieve platform expense easy convenient personalize affordable might mvp min
        imum viable product advise app work customer centre response time gui interfa
        ce app merchant rider also monitor hopefully really work
        69      easy book ride
        153     easy reliable
        207     easy-going
        271     app content interest easy use look forward ofood deal
        278     hello app good easy use think app say ride price ridiculously increase
        hundred naira continiues like do use app thank
        302     ok easy use
        304     good easy
        312     cool easy
        375     easy ride
        Name: reviews_clean, dtype: object
```

## Predicted Negative N-grams

many of users complained of experiencing technical issues of the app not able to connect to the internet on their mobile phones, also geo-location for oride service not accurate

```
In [92]: n_gram(dfn_pred_bad['reviews_clean'], 4, 15)
```

	gram	count
0	(keep, say, something, go)	4
1	(say, something, go, wrong)	4
2	(set, high, accuracy, location)	4
3	(show, something, go, wrong)	3
4	(please, set, high, accuracy)	3
5	(say, internal, server, error)	3
6	(call, customer, care, line)	3
7	(transfer, money, bank, account)	2
8	(something, go, wrong, check)	2
9	(take, long, time, get)	2
10	(check, connection, try, later)	2
11	(refuse, open, keep, say)	2
12	(open, keep, say, something)	2
13	(something, go, wrong, please)	2
14	(go, wrong, please, check)	2

## More Word Clouds









```
# print wordcloud
wd_title = 'Predicted Negative Reviews Word Cloud'
show_wordcloud(dfn_pred_bad['reviews_clean'], title = wd_title)
```

---

Predicted Negative Reviews Word Cloud

